
Table of Contents

.....	1
Question 1	1
Question 2	4
Question 3	7

```
% Sean Bennett  
% HW6
```

Question 1

```
load('CancerMicroarray.mat')  
  
numTumors = length(G)  
numGenes = length(X)  
  
% There are 2308 different genes sequenced from 83 different tumors.  
  
[coeffPC, scorePC, latentPC] = pca(X);  
  
cumulativeLatent = cumsum(latentPC) / sum(latentPC);  
  
numPC = max(find(cumulativeLatent < 0.95));  
  
% You need 52 or 53 PCs to explain 95% of the data  
  
figure;  
hold on;  
stairs(cumulativeLatent, 'Color', 'r');  
title('Explained Variation by PC');  
refline(0, 0.95);  
line([52 52],[0 1]);  
hold off;  
  
linCValAcc = zeros(400,1);  
quadCValAcc = zeros(400,1);  
  
for i = 1:400  
  
    trainingindices = randsample(1:83, 60);  
    testindices = 1:83;  
    testindices(ismember(testindices,trainingindices)) = [];  
  
    trainingdata = X(trainingindices, :);  
    testdata = X(testindices, :);  
  
    trainingclass = G(trainingindices);  
    testclass = G(testindices);
```

```

[trainingcoeffPC, trainingscorePC, traininglatentPC] = pca(trainingdata);

linClass = fitcdiscr(trainingscorePC(:, 1:10), trainingclass);
quadClass = fitcdiscr(trainingscorePC(:, 1:10), trainingclass, ...
    'DiscrimType', 'pseudoQuadratic');

predictedLinClass = predict(linClass, testdata(:,1:10));
predictedQuadClass = predict(quadClass, testdata(:,1:10));

linCValAcc(i) = mean(predictedLinClass == testclass');
quadCValAcc(i) = mean(predictedQuadClass == testclass');

end;

avgLinearCrossVal = mean(linCValAcc)
avgQuadraticCrossVal = mean(quadCValAcc)

% My mean linear and quadratic cross validated accuracies are 22.9% and
% 26.0%. This seems low, but I can't identify anything incorrect in my
% implementation of the PCA/fitcdiscr functionality so I assume there's
% some sort of logic error in my test/training data set creation in the
% loop, or the test data actually passed to predict. I was a little fuzzy
% on that.
% Not sure, but it at least compiles and produces a number!

```

```

numTumors =

```

```

    83

```

```

numGenes =

```

```

    2308

```

```

avgLinearCrossVal =

```

```

    0.2349

```

```

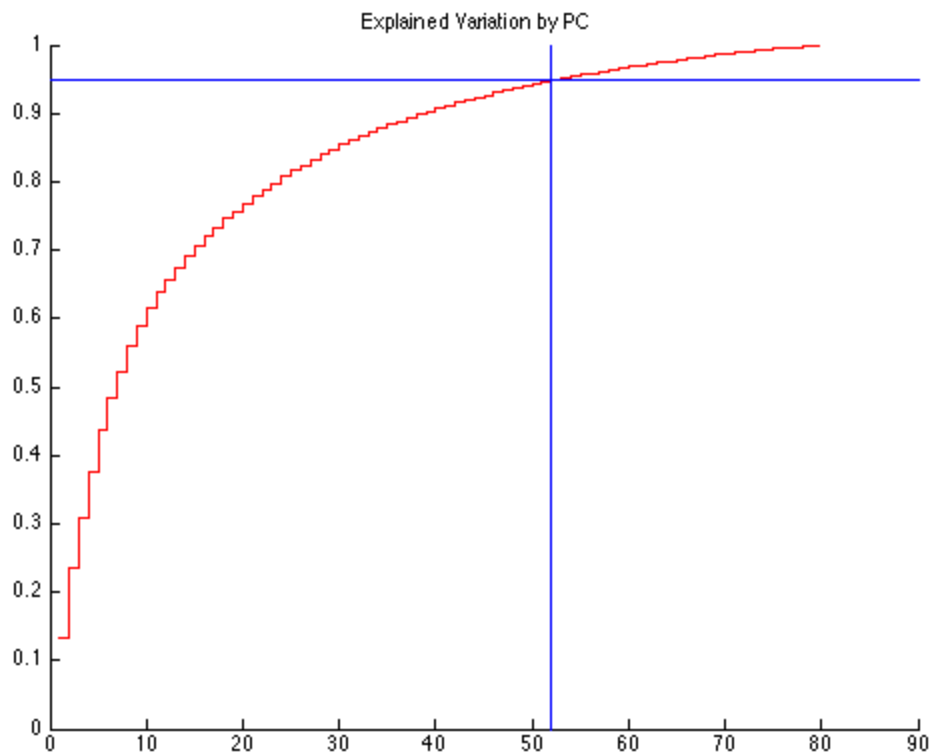
avgQuadraticCrossVal =

```

```

    0.2667

```



```
% Decision trees are super cool. They seek to partition data space into a  
% dichotomous outcome, is a, or is not a for example that minimizes  
% "impurity". It does this recursively, producing some sort of tree that  
% will have multiple branches looking like is a / is not a with is not a  
% subdivided further. Classification trees deal with classifications alone,  
% and then there are regression trees that can deal with continuous data.  
% There are then CART trees, which can deal with both. (I have played with  
% these a little bit in analyzing large ecological data sets where we were  
% trying to determine disease load based upon multiple terrain and  
% oceanographic types)
```

```
treeCVal = zeros(400,1);
```

```
for i = 1:400
```

```
    trainingindices = randsample(1:83, 60);
```

```
    testindices = 1:83;
```

```
    testindices(ismember(testindices,trainingindices)) = [];
```

```
    trainingdata = X(trainingindices, :);
```

```
    testdata = X(testindices, :);
```

```
    trainingclass = G(trainingindices);
```

```
    testclass = G(testindices);
```

```
    [trainingcoeffPC, trainingscorePC, traininglatentPC] = pca(trainingdata);
```

```

treeDiscr = fitctree(trainingscorePC(:, 1:10), trainingclass);

treePredict = predict(treeDiscr, testdata(:, 1:10));

treeCVal(i) = mean(treePredict == testclass');

end
% the Decision tree has ~ 25% cross validated accuracy. Still somewhat
% sucky. I'm pretty sure it's a problem with the predict function, but I
% don't know how to fix it. I'm still unsure of what exactly the first 10
% PCs represents in my data, like how do I access the test data that
% corresponds to those PCs, or is there any way to represent that data?

```

Question 2

```

load('fisheriris');
numSpecies = length(unique(species));
numIris = length(species);
numMeas = size(meas,2);

% The measruements Ronald Fisher took were Petal Width, Petal Length, Sepal
% Width, and Sepal Length, measured in mm There were three different
% species, setsoa, versicolor, and virginica. There were 150 different
% flowers sampled.

[coeffPC2, scorePC2, latentPC2] = pca(meas);

figure;
gscatter(scorePC2(:, 1), scorePC2(:, 2), species);
title('Iris measurements PCs 1 and 2 grouped by actual species');

kmeans2 = kmeans(meas, 2);
figure;
gscatter(scorePC2(:, 1), scorePC2(:, 2), kmeans2);
title('Iris measurements PCs 1 and 2 grouped by K-means = 2');


kmeans3 = kmeans(meas, 3);
figure;
gscatter(scorePC2(:, 1), scorePC2(:, 2), kmeans3);
title('Iris measurements PCs 1 and 2 grouped by K-means = 3');



kmeans4 = kmeans(meas, 4);
figure;
gscatter(scorePC2(:, 1), scorePC2(:, 2), kmeans4);
title('Iris measurements PCs 1 and 2 grouped by K-means = 4');


% Determining the number of groups, k, is hard algorithmically. Some guys
% out of Stanford apparently have developed a Gap statistic that determines
% the k such that the Gap(k) >= Gap(k + 1) - sim error(k+1), It seems to
% work well. Some other guys ahve figured out an algorithm based on a
% statistical test that determines a K in which the data about each k
% centroid closest approximates a Gaussian distribution. There's also a way

```

% to do it graphically via silhouette graphs or information theory
% critirions, but i'm well out of sentence now so lets just say those
% exist.







Question 3

% Pt a.

% DBSCAN uses three different types of points, core points, border points, and noise points to determine clusters based upon the comparison of high density and low density regions. Core points have some minimum number of points within some radius epsilon, border points fall within a neighborhood of a core point, but do not have the minimum number of points to qualify as a core point, and noise points are any points which are not noise or border points. The general algorithm is to classify points as one of the three, remove noise points, and then cluster based upon remaining points.

% Pt b.

```
load('mydata');
```

```
epsilon=0.5;
```

```
MinPts=10;
```

```
IDX=DBSCAN(X,epsilon,MinPts);
```

```
figure;
```

```
PlotClusterinResult(X, IDX);
```

```
title(['DBSCAN Clustering (\epsilon = ' num2str(epsilon) ...
```

```

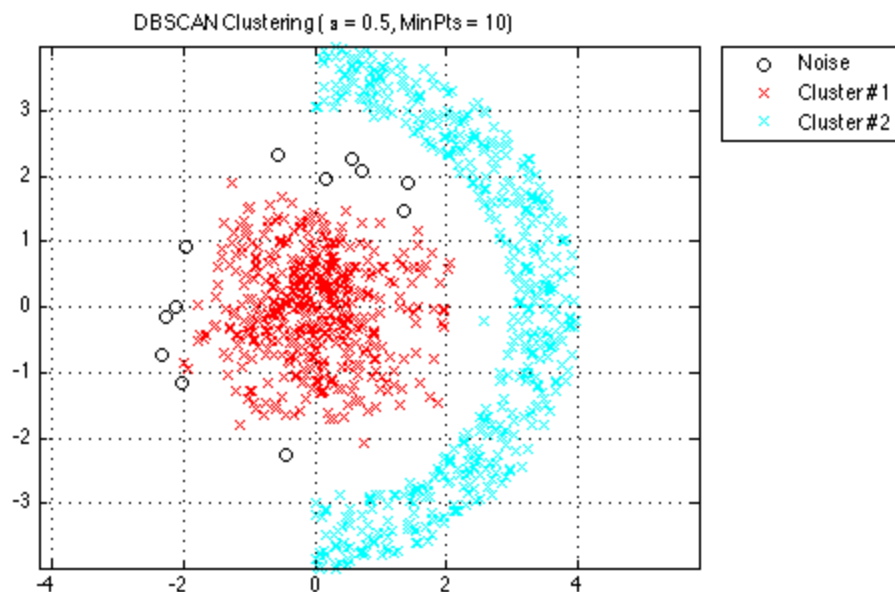
    ', MinPts = ' num2str(MinPts) ')']);


Q3kmeans = kmeans(X, 2);
figure;
gscatter(X(:, 1), X(:, 2), Q3kmeans);

% pt c.

% The DBSCAN algorithm is density based and has the ability to deal with
% non-linear groupings. Visible in the comparison between the two graphs is
% k-means needs to draw a straight line between the two groups, while
% DBSCAN can draw a non-straight/curved line. I found a thing that
% describes it as k-means can't deal with "non-globular" shapes.

```





Published with MATLAB® R2014a